

# MA-INF 4223 - Lab Distributed Big Data Analytics

## Spark Fundamentals II

Dr. Hajira Jabeen, Gezim Sejdiu

Winter Semester 2018-19



# Lesson objectives

- ❖ After completing this lesson, you should be able to:
  - Understand and use various Spark Libraries
    - Spark SQL
    - Spark GraphX - graph processing



# Pair RDDs

- A common form of data processing
- Main intuition behind the mapreduce
- Often beneficial to project down the complex data types to Key-value pairs

Distributed key-value pairs

- Additional specialised methods for working with data associated with Keys
- `groupByKey()`, `reduceByKey()`, `join`



# Pair RDDs

Creation: Mostly from existing non-pair RDDs

E.g.

```
val pairRdd = rdd.map(page => (page.title, page.text))
```

- groupByKey
- reduceByKey
- mapValues
- keys
- Join
- leftOuterJoin/rightOuterJoin
- countByKey



# Shuffling

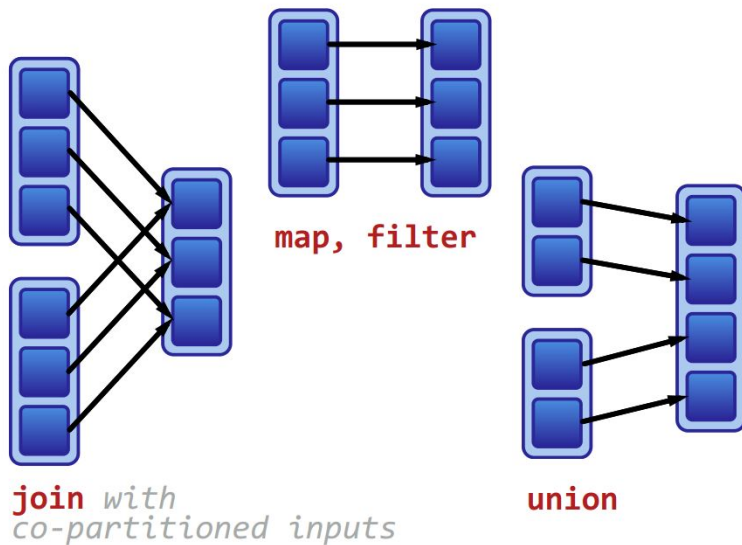
- `groupByKey()`
  - Shuffles all the keys across network to combine all the keys
- `reduceByKey(func: (V, V) => V): RDD[(K, V)]`
  - Conceptually, `reduceByKey` can be thought of as a combination of first doing `groupByKey` and then reducing on all the values grouped per key.
  - Reduces on the mapper side first
  - Reduce again after shuffling
  - Less data needs to be sent over the network
  - Non trivial gains in performance



# Dependencies / Shuffling

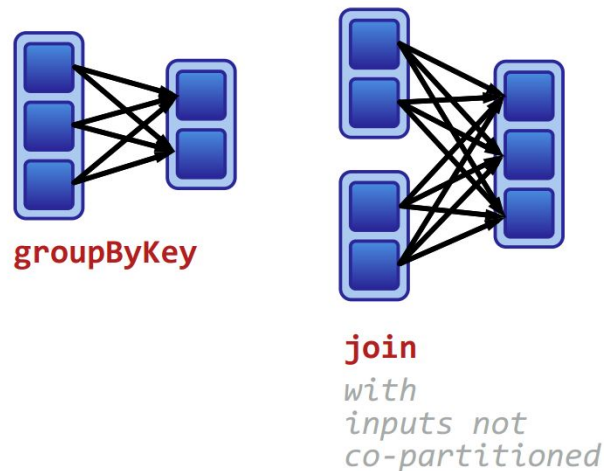
## Narrow dependencies:

Each partition of the parent RDD is used by at most one partition of the child RDD.



## Wide dependencies:

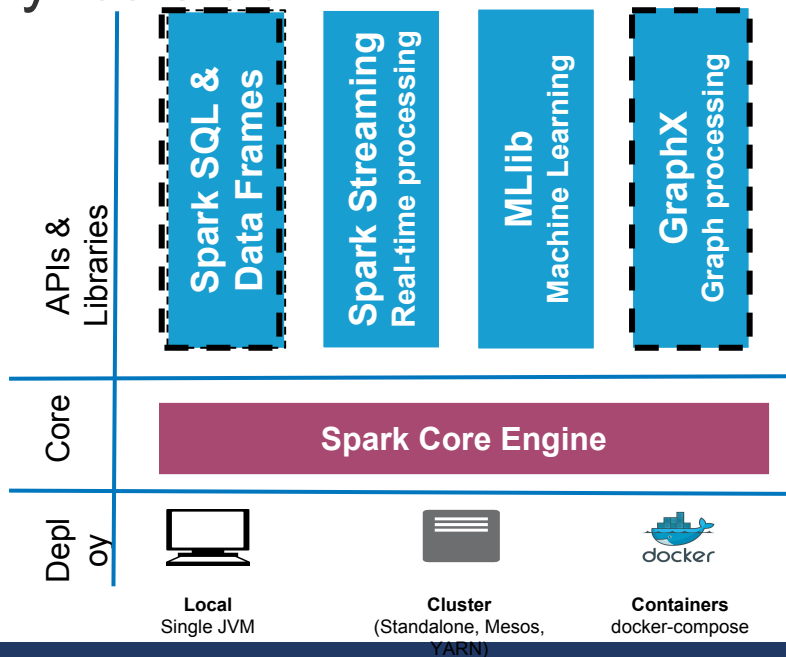
Each partition of the parent RDD may be depended on by multiple child partitions.





# Spark Libraries

## ❖ A unified analytics stack





# Overview

- ❖ [Spark SQL: Relational Data Processing in Spark](#)
- ❖ [GraphX: A Resilient Distributed Graph System on Spark](#)



# Spark SQL



# Motivation

- ❖ Support relational processing both within Spark programs
- ❖ Provide high performance with established DBMS techniques
- ❖ Easily support new data sources, including semi-structured data and external databases amenable to query federation
- ❖ Enable extension with advanced analytics algorithms such as graph processing and machine learning



# Motivation

- ❖ Users:
  - Want to perform ETL-relational processing
    - data Frames
  - Analytics - procedural tasks
    - UDFs



# Spark SQL

- ❖ A module that integrates relational processing with Spark's Functional programming API
- ❖ Spark SQL allows relational processing
- ❖ Perform complex analytics
  - Integration between relational and procedural processing through declarative Data Frame
  - Optimizer ( catalyst)
    - Composable rules
    - Control code generation
    - Extension points
    - Schema inference for json
    - ML types
    - Query federation



# Spark SQL

Three main APIs

- SQL Syntax
- DataFrames
- Datasets

Two specialised backend components

- Catalyst
- Tungsten



# Data Frame

- ❖ DataFrames are collections of structured records that can be manipulated using Spark's procedural API,
- ❖ Supports relational APIs that allow richer optimizations.
- ❖ Created directly from Spark's built-in distributed collections of Java/Python objects,
- ❖ Enables relational processing in existing Spark Programs
- ❖ DataFrame operations in SparkSQL go through a relational optimizer, Catalyst



# Catalyst

- ❖ Catalyst is the first production quality query optimizer built on such functional language.
- ❖ It contains an extensible query optimizer
- ❖ Catalyst uses features of the Scala programming language,
  - Pattern-matching
  - Express composable rules
  - Turing complete language

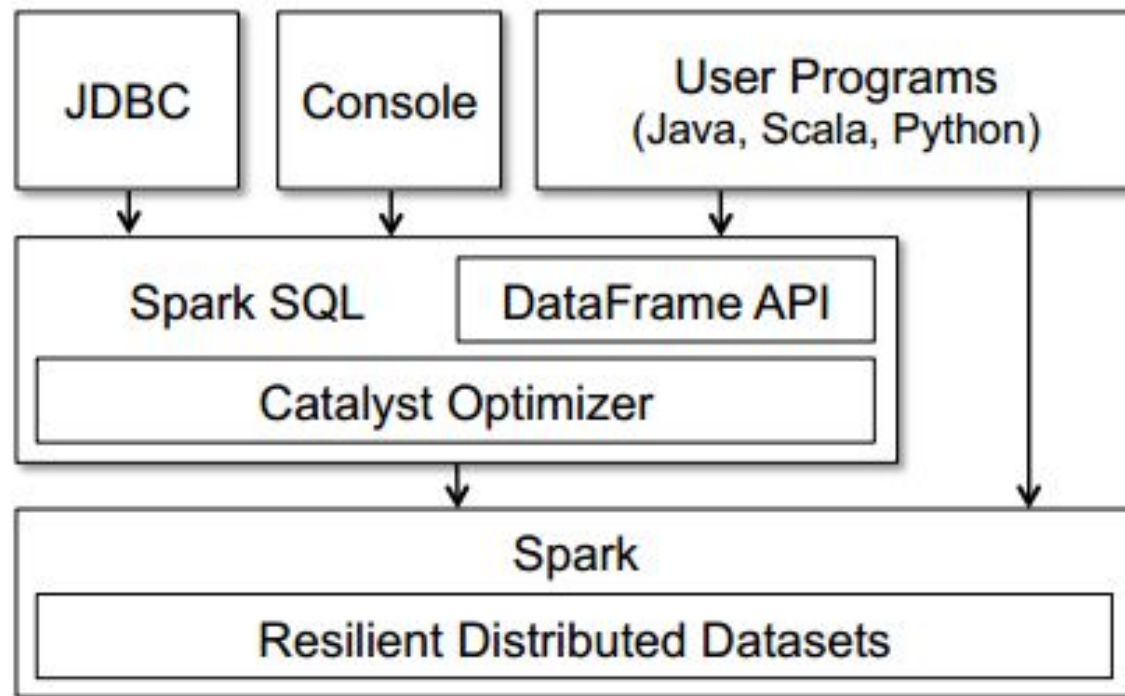


# Catalyst

- ❖ Catalyst can also be
  - extended with new data sources,
  - semi-structured data
    - such as JSON
    - “smart” data stores to use push filters
    - e.g., HBase
    - user-defined functions;
    - User-defined types for domains e.g. machine learning.
- ❖ Spark SQL simultaneously makes Spark accessible to more users and improves optimizations



# Spark SQL





# DataFrame

- ❖ DataFrame is a distributed collection of rows with the “Known” schema like table in a relational database.
- ❖ Each DataFrame can also be viewed as an RDD of Row objects, allowing users to call procedural Spark APIs such as map.
- ❖ Spark DataFrames are lazy, in that each DataFrame object represents a logical plan to compute a dataset, but no execution occurs until the user calls a special “output operation” such as save



# DataFrame

- ❖ Created from an RDD using `.toDF()`
- ❖ Reading from a file `()`



# Example

- `ctx = new HiveContext()`
- `users=ctx.table("users")`
- `young = users.where(users("age")<21)`
- `println(young.count())`



# Data Model

- ❖ DataFrames support all common relational operators, including
  - projection (select),
  - filter (where),
  - join, and
  - aggregations (groupBy).
- ❖ Users can break up their code into Scala, Java or Python functions that pass DataFrames between them to build a logical plan, and will still benefit from optimizations across the whole plan when they run an output operation.



# Optimization

- ❖ The API analyze logical plans eagerly
  - identify whether the column names used in expressions exist in the underlying tables,
  - whether the data types are appropriate
- ❖ Spark SQL allows users to construct DataFrames directly against RDDs of objects native to the programming language.
- ❖ Spark SQL can automatically infer the schema of these objects using reflection



# Optimization

- ❖ Uses columnar cache
  - reduce memory footprint by an order of magnitude because it applies columnar compression schemes such as dictionary encoding and run-length encoding.
- ❖ In Spark SQL, UDFs can be registered inline by passing Scala, Java or Python functions, which may use the full Spark API internally.



# Catalyst - Extension

- ❖ Easy to add new optimization techniques and features to Spark SQL
- ❖ Enable external developers to extend the optimizer
  - E.g. adding data source specific rules that can push filtering or aggregation into external storage systems,
  - support for new data types.
- ❖ Catalyst supports both rule-based and cost-based optimizations





# Catalyst

- ❖ Catalyst contains a general library for representing trees(Abstract Syntax Tree) and applying rules to manipulate them
- ❖ Catalyst offers several public extension points, including external data sources and user-defined types.



# Trees

- ❖ The main data type in Catalyst is a tree composed of node objects.
- ❖ Each node has a node type and zero or more children.
- ❖ New node types are defined in Scala as subclasses of the `TreeNode` class.
- ❖ These objects are immutable and can be manipulated using functional transformations



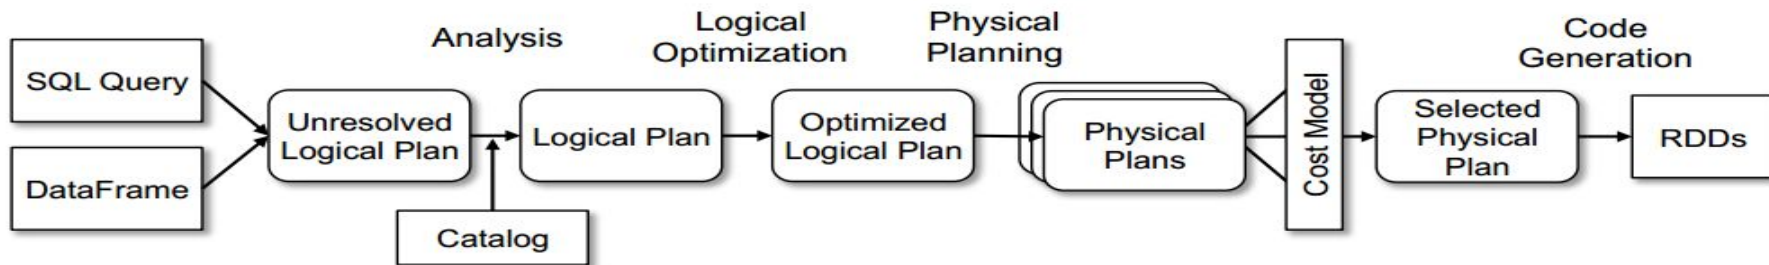
# Rules

- ❖ Trees can be manipulated using rules, which are functions from a tree to another tree.
- ❖ While a rule can run arbitrary code on its input tree (given that this tree is just a Scala object),
- ❖ the most common approach is to use a set of pattern matching functions that find and replace subtrees with a specific structure.



# Tree Transformation

- ❖ Catalyst's general tree transformation framework works in four phases
  - analyzing a logical plan to resolve references
  - logical plan optimization
  - physical planning
  - code generation to compile parts of the query to Java bytecode.



# Spark GraphX



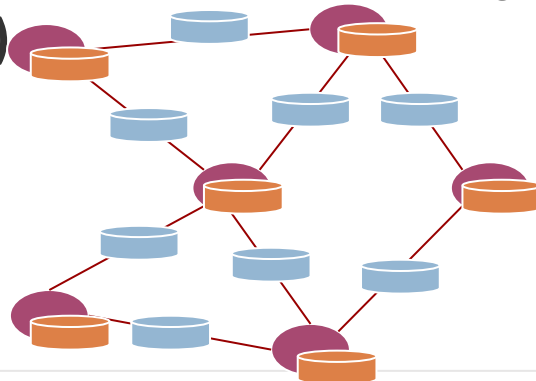
# Spark GraphX

- ❖ Graph computation system which runs in the Spark data-parallel framework.
- ❖ GraphX extends Spark's Resilient Distributed Dataset (RDD) abstraction to introduce the Resilient Distributed Graph (RDG)



# Spark GraphX

- ❖ Spark GraphX - stands for graph processing
  - For graph and graph-parallel computation
- ❖ At a high level, GraphX extends the Spark RDD by introducing a new Graph abstraction:
  - a directed multigraph with properties attached to each vertex and edge.
- ❖ It is based on Property Graph model →  $G(V, E)$ 
  - Vertex Property
    - Triple details
  - Edge Property
    - Relations
    - Weights





# Resilient Distributed Graph (RDG)

- ❖ A tabular representation of the efficient vertex-cut partitioning and data-parallel partitioning heuristics
- ❖ Supports implementations of the
  - PowerGraph and
  - Pregel graph-parallel
- ❖ Preliminary performance comparisons between a popular dataparallel and graph-parallel frameworks running PageRank on a large real-world graph





# Graph Parallel

- ❖ Graph-parallel computation typically adopts a vertex (and occasionally edge) centric view of computation
- ❖ Retaining the **data-parallel metaphor**, program logic in the GraphX system defines transformations on graphs with each operation yielding a new graph
- ❖ The core data-structure in the GraphX systems is an immutable graph



# GraphX operations

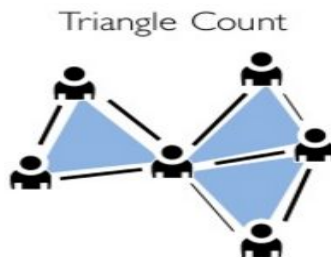
```
class Graph[VD, ED] {  
  // Information about the Graph  
  val numEdges: Long  
  val numVertices: Long  
  val inDegrees: VertexRDD[Int]  
  val outDegrees: VertexRDD[Int]  
  val degrees: VertexRDD[Int]  
  
  // Views of the graph as collections  
  val vertices: VertexRDD[VD]  
  val edges: EdgeRDD[ED]  
  val triplets: RDD[EdgeTriplet[VD, ED]]  
  
  // Functions for caching graphs  
  def persist(newLevel: StorageLevel = StorageLevel.MEMORY_ONLY): Graph[VD, ED]  
  def cache(): Graph[VD, ED]  
  def unpersistVertices(blocking: Boolean = true): Graph[VD, ED]  
  // Change the partitioning heuristic  
  def partitionBy(partitionStrategy: PartitionStrategy): Graph[VD, ED]  
  // Transform vertex and edge attributes  
  def mapVertices[VD2](map: (VertexId, VD) => VD2): Graph[VD2, ED]  
  def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]  
  ----
```



# GraphX build-in Graph Algorithms

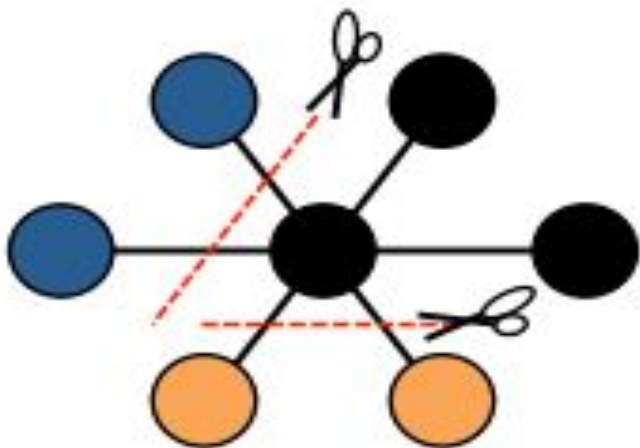
// Basic graph algorithms

```
=====
def pageRank(tol: Double, resetProb: Double = 0.15): Graph[Double, Double]
def connectedComponents(): Graph[VertexId, ED]
def triangleCount(): Graph[Int, ED]
def stronglyConnectedComponents(numIter: Int): Graph[VertexId, ED]
```

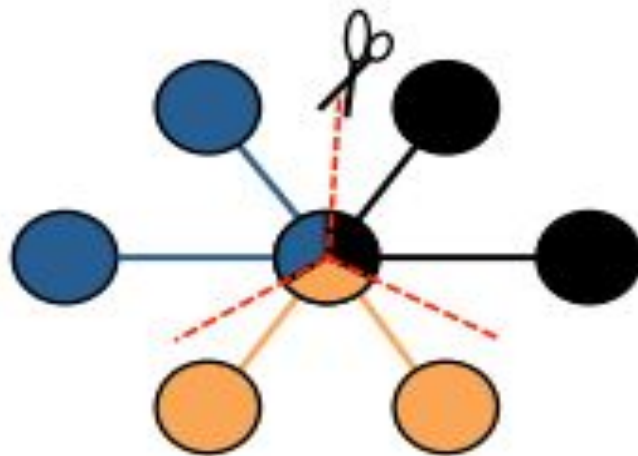




# Edge-Cut vs Vertex-Cut



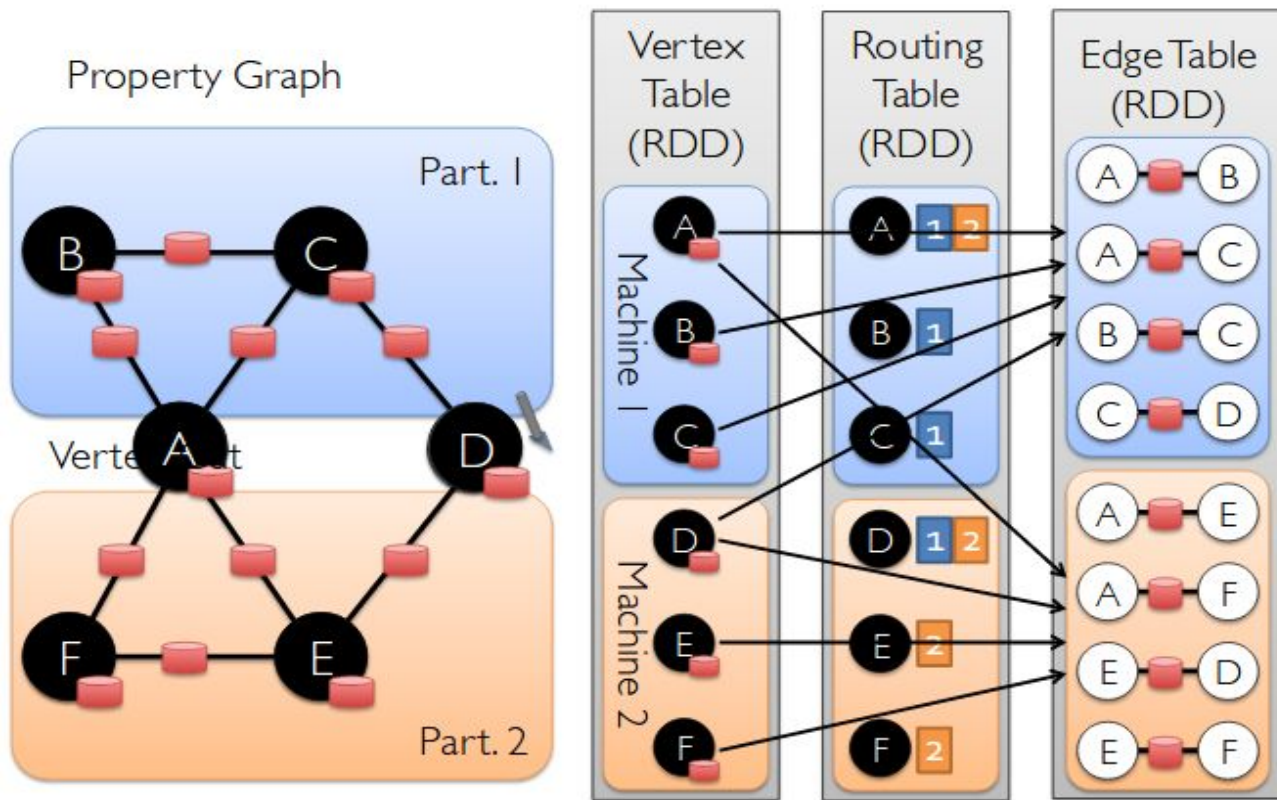
(a) Edge-Cut



(b) Vertex-Cut



# Encoding Property Graphs as RDDs





# Edge Table

- ❖ EdgeTable(pid, src, dst, data): stores the adjacency structure and edge data
- ❖ Each edge is represented as a tuple consisting of the
  - source vertex id,
  - destination vertex id,
  - user-defined data
  - virtual partition identifier (pid).



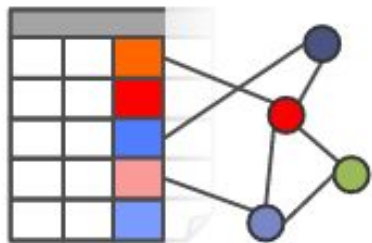
# Vertex Data Table

- ❖ `VertexDataTable(id, data)`: stores the vertex data, in the form of a vertex (id, data) pairs
- ❖ `VertexMap(id, pid)`: provides a mapping from the id of a vertex to the ids of the virtual partitions that contain adjacent edges



## New API

*Blurs the distinction between  
Tables and Graphs*



## New Library

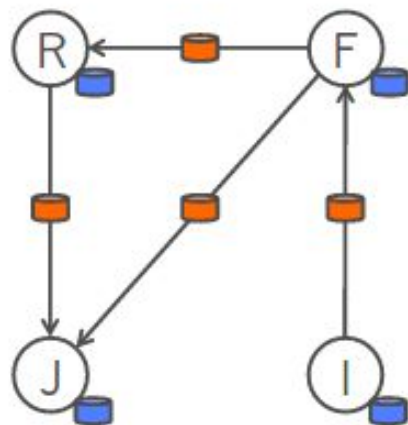
*Embeds Graph-Parallel  
model in Spark*







## Property Graph



## Vertex Table

Id	Attribute (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk.)
Istoica	(Prof., Berk.)

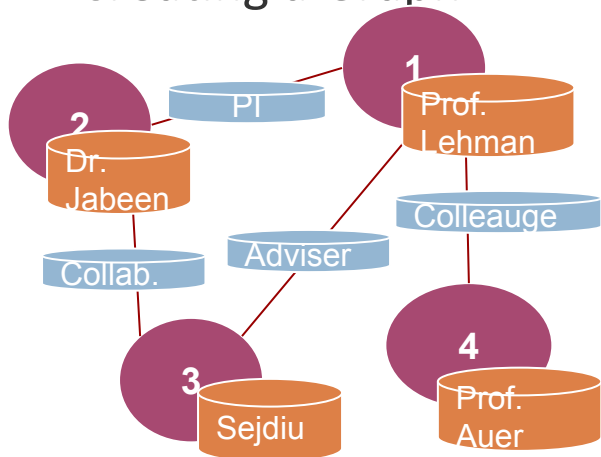
## Edge Table

SrcId	DstId	Attribute (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI



# Spark GraphX - Getting Started

## ❖ Creating a Graph



```
type VertexId = Long
// Create an RDD for the vertices
val users: RDD[(VertexId, (String, String))] =
    spark.sparkContext.parallelize(
        Array((3L, ("sejdiu", "phd_student")),
              (2L, ("jabeen", "postdoc")),
              (1L, ("lehmann", "prof")),
              (4L, ("auer", "prof"))))

// Create an RDD for edges
val relationships: RDD[Edge[String]] =
    spark.sparkContext.parallelize(
        Array(Edge(3L, 2L, "collab"),
              Edge(1L, 3L, "advisor"),
              Edge(1L, 4L, "colleague"),
              Edge(2L, 1L, "pi")))

// Build the initial Graph
val graph = Graph(users, relationships)
```

Vertex RDD	
vID	Property(V)
1L	(lehmann, prof)
2L	(jabenn, postdoc)
3L	(sejdiu, phd_student)
4L	(auer, prof)

Edge RDD		
sID	dID	Property(E)
1L	3L	advisor
1L	4L	colleague
2L	1L	pi
3L	2L	collab



# GraphX Optimizations

- ❖ Mirror Vertices
- ❖ Partial materialization
- ❖ Incremental view
- ❖ Index Scanning for Active Sets
- ❖ Local Vertex and Edge Indices
- ❖ Index and Routing Table Reuse



# References

- [1]. [Spark SQL: Relational Data Processing in Spark](#) by Armbrust, Michael, Reynold Xin, Cheng Lian, Yin Huai, Davies Liu, Joseph K. Bradley, Xiangrui Meng, Tomer Kaftan, Michael J. Franklin, Ali Ghodsi and Matei Zaharia *in SIGMOD Conference*, 2015.
- [2]. “Spark SQL, DataFrames and Datasets Guide” - <http://spark.apache.org/docs/latest/sql-programming-guide.html>
- [3]. [GraphX: Graph Processing in a Distributed Dataflow Framework](#) by Gonzalez, Joseph, Reynold Xin, Ankur Dave, Daniel Crankshaw, Michael J. Franklin and Ion Stoica *in OSDI*, 2014.
- [4]. “GraphX Programming Guide” - <http://spark.apache.org/docs/latest/graphx-programming-guide.html>

# THANK YOU !

<http://sda.cs.uni-bonn.de/teaching/dbda/>

- <http://sda.cs.uni-bonn.de/>
- <https://github.com/SANSA-Stack>
- <https://github.com/big-data-europe>
- <https://github.com/SmartDataAnalytics>



**Dr. Hajira Jabeen**

[jabeen@cs.uni-bonn.de](mailto:jabeen@cs.uni-bonn.de)

Room 1.066 (Appointment per e-mail)



**Gezim Sejdiu**

[sejdiu@cs.uni-bonn.de](mailto:sejdiu@cs.uni-bonn.de)

Room 1.068 (Appointment per e-mail)