UNIVERSITÄT BONN

SMART DATA ANALYTICS
FROM DATA TO KNOWLEDGE

MA-INF 4223 - Lab Distributed Big Data Analytics

Spark Fundamentals II

Dr. Hajira Jabeen, Gezim Sejdiu

Winter Semester 2018-19

Lesson objectives

- After completing this lesson, you should be able to:
  - Understand and use
    - Spark MLlib

Spark ML

Overview

- MLlib: Machine Learning in Apache Spark

Spark ML

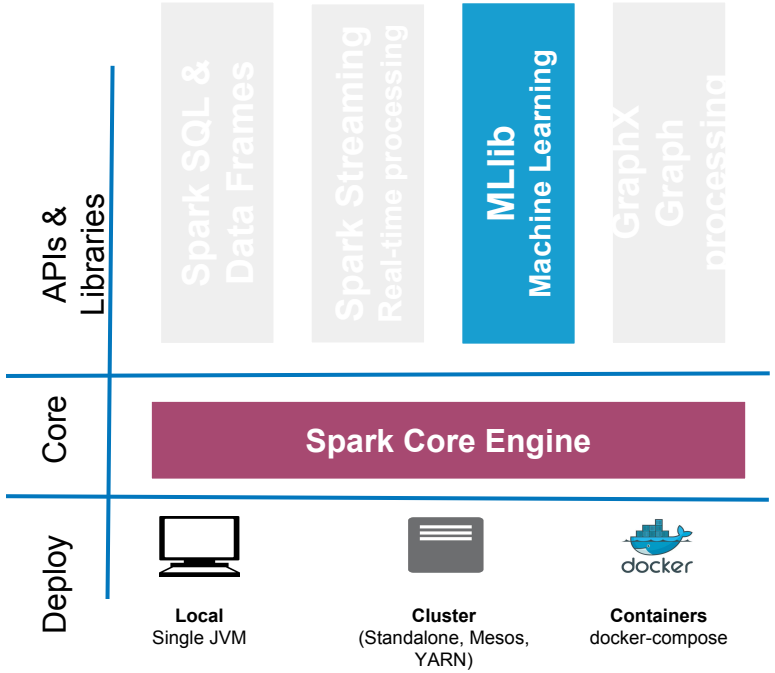
APIs & Libraries
Spark SQL & Data Frames
Spark Streaming Real-time processing
MLlib Machine Learning
GraphX Graph processing
Core
Spark Core Engine
Deploy
Local
Single JVM
Cluster
(Standalone, Mesos, YARN)
Containers
docker-compose
docker

ML Algorithms overview

- Machine learning are separated in two major types of algorithms :
  - Supervised - labeled data in which both, input and output are provided to the algorithm
  - Unsupervised - do not have the outputs in advance

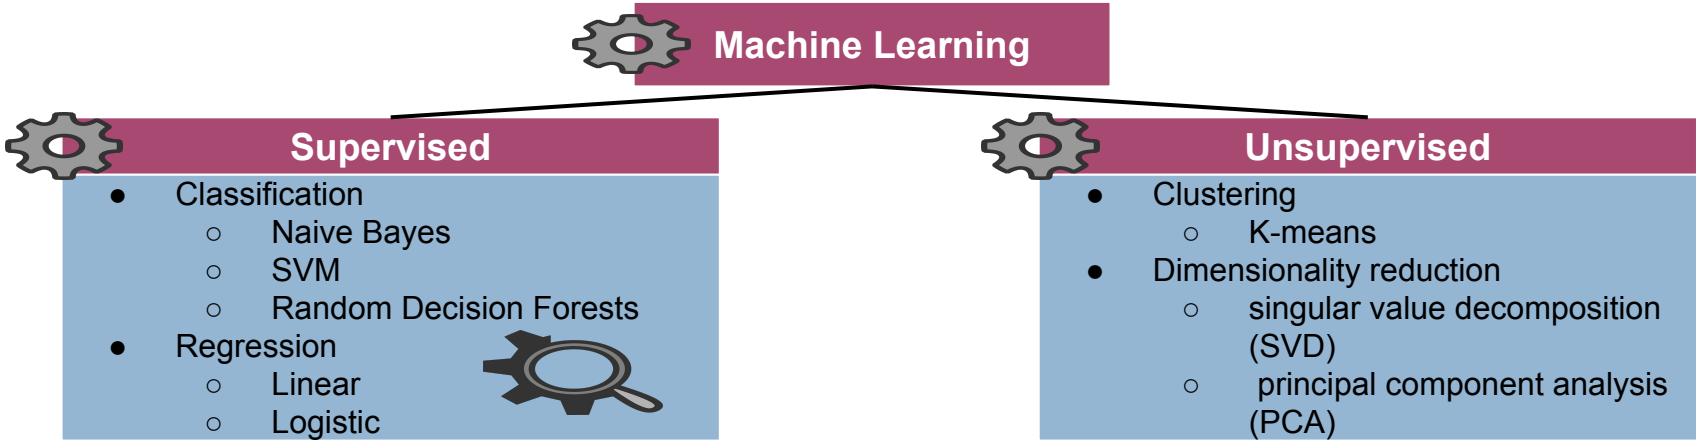

Spark ML

- MLlib is a standard component of Spark providing machine learning primitives on top of Spark
- It is scalable machine learning, statistics, math libraries
- Supports out-of-the-box most popular machine learning algorithms like Linear regression, Logistic regression, Decision Trees
- Is available in Scala, Java, Python, and R

Spark ML-pipelines

- Uniform set of APIs for creating and tuning data processing/machine learning pipelines
- Core concepts:
  - DataFrame: RDD with named columns. SQL-like syntax and other core RDD operations
  - Transformer: DataFrame => DataFrame. Eg., features to predictions(classifier)
  - Estimator: DataFrame => Transformer. e.g., learning algorithm
  - Param: map of params
  - Pipeline: Chain of Transformers and Estimators. Specifies the data flow

Spark ML-pipelines

- Transformer
  - A Transformer is an abstraction which uses an algorithm to transform one DataFrame to another
  - It implements a method `transform()`

Spark ML-pipelines

- Estimator
  - An Estimator abstraction uses an algorithm which is fitted into a DataFrame returning a model
  - It implements a method `fit()`

Spark ML-pipelines Example

- Split text into words => convert numerical features => generate a prediction model

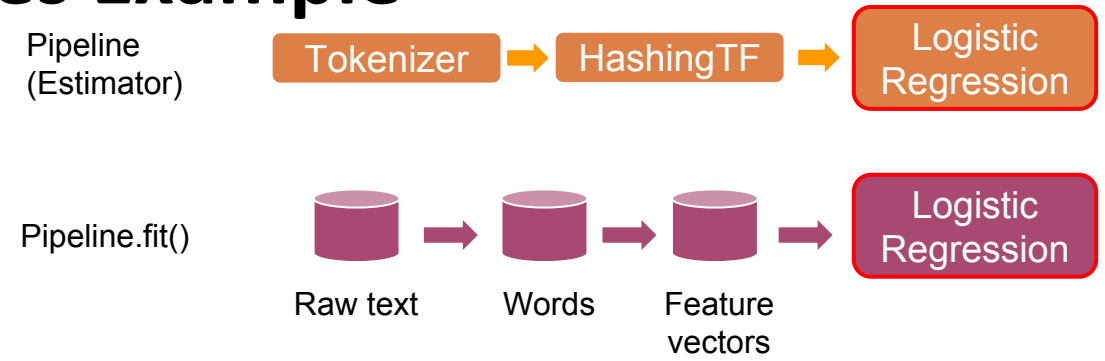


```
val tokenizer = new Tokenizer().setInputCol("text").setOutputCol("words")
val hashingTF = new HashingTF().setNumFeatures(1000).setInputCol(tokenizer.getOutputCol)
.setOutputCol("features")
val lr = new LogisticRegression().setMaxIter(10).setRegParam(0.01)
val pipeline = new Pipeline().setStages(Array(tokenizer, hashingTF, lr))
val model = pipeline.fit(training.toDF)
val test = sc.parallelize(Seq(
Document(4L, "spark i j k"),
Document(5L, "l m n"),
Document(6L, "mapreduce spark"),
Document(7L, "apache hadoop")))
val predictions = model.transform(test.toDF)
```

References

[1]. MLlib: Machine Learning in Apache Spark by Meng, Xiangrui, Joseph K. Bradley, Burak Yavuz, Evan R. Sparks, Shivaram Venkataraman, Davies Liu, Jeremy Freeman, D. B. Tsai, Manish Amde, Sean Owen, Doris Xin, Reynold Xin, Michael J. Franklin, Reza Bosagh Zadeh, Matei Zaharia and Ameet Talwalkar *in Journal of Machine Learning Research 17*, 2016.

[2]. “Machine Learning Library (MLlib) Guide” - http://spark.apache.org/docs/latest/ml-guide.html

THANK YOU !

http://sda.cs.uni-bonn.de/teaching/dbda/

- http://sda.cs.uni-bonn.de/
- https://github.com/SANSA-Stack
- https://github.com/big-data-europe
- https://github.com/SmartDataAnalytics

Dr. Hajira Jabeen
jabeen@cs.uni-bonn.de
Room 1.066 (Appointment per e-mail)

Gezim Sejdiu
sejdiu@cs.uni-bonn.de
Room 1.068 (Appointment per e-mail)